

YUANPEI CAO

215-720-6285 | yuanpei.cao@airbnb.com

WORKING EXPERIENCE

Airbnb - Data Scientist

Apr. 2018 -

- Smart Routing: lead end-to-end statistics project for optimizing payment routes between multiple processing providers
 - Developed heuristic tree segmentation algorithm (using statistical inference and offline reply estimate) to find the opportunity route segments.
 - Optimized the first attempt acceptance rate without violating contractual obligations from different processor companies (Chase, WorldPay, etc).
 - Increased first attempt acceptance rate by x% compared to current system according to offline reply estimate, which can be translated to 2-3 millions dollars annual revenue increase.
 - Published internal knowledge post and submitted Airbnb invention disclosure form to legal team.
- Provided analytics support (anomaly detection / new payment service monitoring / data pipeline & dashboard) for multiple high-visibility strategic partnerships, resulting in millions dollars in bottom-line savings.
- Designed and analyzed A/B tests for Guest & Host Commerce (Pay Less Upfront, Alipay Payout, Quickpay, etc).

AdColony (Opera Software) - Data Scientist

Jan. 2017 - Mar. 2018

Ad Install Rate Prediction: lead end-to-end machine learning project for IR prediction on new ads (cold start)

- Handled with raw data cleaning and feature engineering from more than 100TB historical data sources, including the query, the text and video of the ad creative, and various ad-related metadata.
- Developed the ensemble learning framework (k-NN, gradient boosting regression tree, a variety of content-based similarity models using deep learning and NLP techniques, etc) to find the target users.
- Automated pipeline using disparate tools/sources like Shell, Python, BigQuery, Redshift, s3, ec2, Cron.
- Increased install rate and eCPM by more than 300% compared to previous models.
- Presented the work to CTO and other technical leaders; Also presented to non-technical audience from business team.

Department of Biostatistics, University of Pennsylvania - Postdoc

Sep. 2016 - Jan. 2017

- Developed statistical models (hypothesis testing, multivariate analysis, network estimation, missing value recovery) for high-dimensional compositional data with applications to human microbiome study.
- Handled with raw data cleaning from cross-section study and clinical trials, and implemented the statistical algorithm and data visualization techniques by R/MATLAB/Python/Cytoscape.
- Collaborated with experts from Children Hospital of Philadelphia (CHOP) to explore the relationships between microbiome and human health.

KNOWLEDGE AND SKILLS

Languages: Python (Scipy, Pandas, Scikit-learn, XGBoost, Tensorflow, etc), R, Java, Shell, SQL, Matlab
Cloud: AWS (ec2, S3, EBS, Redshift), Google Cloud (BigQuery), MySQL, Presto, Hive, Airflow, Aerospike
Miscellaneous: Tableau, Linux, Cron, scripting automation, git, tmux, LaTeX, experienced in Hadoop, Apach Spark

Quantitative Analysis: Ph.D. level knowledge in statistics, working experience in machine learning and A/B testing

EDUCATION

University of Pennsylvania	<i>Sep. 2011 - Aug. 2016</i>
Ph.D. in Applied Mathematics and Computational Science	<i>Overall GPA: 3.96/4.00</i>
The Wharton School, University of Pennsylvania	<i>Sep. 2011 - May 2016</i>
M.A. in Statistics	<i>Overall GPA: 4.00/4.00</i>
Fudan University, Shanghai, China	<i>Sep. 2007 - Jun. 2011</i>
B.S. in Applied Mathematics (with honors)	<i>Overall GPA: 3.72/4.00 (Top 5%)</i>

PREPRINTS AND WORKING PAPERS

-
- **Cao, Y.**, Lin, W. and Li, H. (2018): *Large Covariance Estimation for Compositional Data via Composition-Adjusted Thresholding*, *Journal of the American Statistical Association (JASA)*.
 - **Cao, Y.**, Lin, W. and Li, H. (2018): *Two-sample Mean Tests for High Dimensional Compositional Data*, *Biometrika*, 105(1): 115-132.
 - **Cao, Y.**, Zhang, A. and Li, H.: *Multi-sample Estimation of Bacterial Composition Matrix in Metagenomics Data*, submitted to *Biometrika*, under second-round review, arXiv:1706.02380.
 - Ma, R., **Cao, Y.**, and Li, H.: *Sparse High-Dimensional Precision Matrix Estimation for Compositional Data*, submitted.
 - Feng, X., Wang, S., Gao, S., **Cao, Y.**, and Murray, A.T.: *MOTO: A Multi-Objective Trajectory Optimization Method for Finding Sequential Activity Locations for Multiple Moving Objects along Road Networks*, submitted to *Environment and Planning B: Urban Analytics and City Science*, under revision.
 - Jie, C., **Cao, Y.**: *GARCH Modeling and Extreme Value Theory-based Fund Risk Measurement and Performance Evaluation*, *Modern Business* (2011) (in Chinese).

KAGGLE COMPETITION

New York City Taxi Fare Prediction	<i>Sept. 2018</i>
Hosted by Google Cloud and Coursera	<i>Leaderboard Rank: 9/1488 (top 1%)</i>

CERTIFICATION

Deep Learning Nanodegree on Udacity	<i>Oct. 2018</i>
Convolutional Neural Networks by deeplearning.ai on Coursera	<i>Jan. 2018</i>
Structuring Machine Learning Projects by deeplearning.ai on Coursera	<i>Jan. 2018</i>
Improving Deep Neural Networks: Hyperparameter tuning, Regularization and Optimization by deeplearning.ai on Coursera	<i>Jan. 2018</i>
Neural Networks and Deep Learning by deeplearning.ai on Coursera	<i>Jan. 2018</i>
Shell Scripting: Discover How to Automate Command Line Tasks on Udemy	<i>May 2017</i>
Hadoop Platform and Application Framework on Coursera	<i>Sep. 2016</i>
Introduction to Apache Spark on EdX	<i>Aug. 2016</i>